





# **APEX: Governed Enterprise AI Gateway**

Safely enabling enterprise AI capabilities through policy, identity, audit, and zero-trust enforcement.

# The Problem



## Uncontrolled AI Capability Access Is Becoming Enterprise Risk



**Note:** Enterprise risk problem exists for MCP, AI agents, APIs, copilots, workflows, and AI tooling

### DIRECT PROMPT INJECTION: ATTACKER MODIFIES USER INPUT

Malicious instructions are **directly** embedded by the attacker into the **user input** sent to the LLM application.



**References:** TechCrunch (May 2023), CVE-2025-68143/68144/68145 (Cyata, Jan 2026), EU AI Act Article 99 (Regulation 2024/1689), Cisco State of AI Security 2026

# Real-World Examples of the Problem Set

## Uncontrolled AI capability access is now enterprise risk

Two incidents bookend the change. One was an employee mistake. The other was the protocol itself.

### INCIDENT 1 · MARCH 2023

samsung

#### Samsung leaks source code into ChatGPT

Three separate incidents in 20 days at Samsung's semiconductor division. Engineers pasted source code, meeting transcripts, and equipment defect identification code into ChatGPT to debug, optimize, and summarize.

##### What it cost

- Proprietary semiconductor code in third-party training data
- Internal meeting content outside the firewall
- Company-wide ban on generative AI (April 2023)
- Reputational and competitive exposure that is irreversible

*The mechanism was an employee using a public tool. No malice. No vulnerability. Just unrestricted capability access.*

### INCIDENT 2 · JANUARY 2026

mcp-server-git

#### Three CVEs in Anthropic's own MCP server

CVE-2025-68143, 68144, 68145. Cyata researchers showed that a malicious README, GitHub issue, or webpage could trigger code execution, file deletion, or context injection. No credentials required.

##### What it means

- The protocol creator's own server had injection surface
- Indirect prompt injection bypasses identity entirely
- Affected every release before December 8, 2025
- The same class of flaw exists across the MCP ecosystem

*The mechanism was the protocol itself. The vendor wrote it. The vendor missed it. Every adopter inherits the risk.*

### The pattern between them

In 2023, an employee was the attack surface. In 2026, the protocol is. Both share one root cause: enterprises gave AI access to enterprise capabilities without a governance layer to mediate it.

**85%+**

Adaptive prompt-injection success against tested coding agents

**13.4%**

Of audited agent skills had critical-level security flaws

**29%**

Of enterprises ready to secure agentic AI deployments

Sources: Bloomberg (May 2023), TechCrunch (May 2023). Cyata / Infosecurity Magazine (Jan 2026): CVE-2025-68143, CVE-2025-68144, CVE-2025-68145. Cisco State of AI Security 2026. ClawHub agent skills audit (2026).

# THE RISING WAVE OF AI RISKS

MCP SERVERS, AGENTS,  
AND EMERGING THREATS

UNRESTRICTED  
MCP SERVERS

DATA EXFILTRATION  
& LEAKAGE

PROMPT  
INJECTION

OVERLY BROAD  
API ACCESS

IT INFRA  
INFOSEC  
TEAMS

## Why NOW?

### Four pressures converging in 2026

AI capability access went from research curiosity to enterprise-wide risk in eighteen months.



#### MCP proliferation

Nov 2024 → today

Model Context Protocol went from Anthropic launch to industry standard in 18 months. Microsoft, OpenAI, Google, AWS, and dozens of dev tools now support it.



#### Documented incidents are now monthly

Q1 2026

Three CVEs in Anthropic's own Git MCP server (Jan 2026). Gemini hijacked through calendar invitations. Microsoft Copilot session hijacking via Reprompt attack. 85%+ adaptive prompt-injection success rates across all tested coding agents.



#### EU AI Act enforcement begins

Aug 2, 2026

High-risk AI systems become enforceable. Penalties up to €35M or 7% of global revenue. Documentation, audit, and human oversight are not optional. Applies to any AI deployed by, or serving, EU operations.



#### Boards are asking, CISOs are unprepared

Cisco 2026

Cisco State of AI Security 2026: most enterprises plan to deploy agentic AI, but only 29% report being prepared to secure those deployments. The gap between board-level ambition and security readiness is the largest in a decade.


#### The window

Capability is here. Incidents are public. Regulation is dated. Board pressure is real. Enterprises will buy governance in 2026, and the question is whether they buy it before or after their first incident.

# The APEX Architecture

## Governed MCP Gateway — Request Flow

AI never holds credentials. Every call passes through identity, policy, and audit.




Plus External Resources: Office, Outlook, Web, etc.

# Governance Layers


## What Each Layer Actually Enforces

A single tool call is checked five times before it touches data.



### Defense in depth, not defense by one prompt rule.

Each layer answers a different question. A jailbroken model can't talk its way past layer 3 (Entra) or layer 4 (RBAC) — those don't read English. And layer 5 means even a successful call leaves a record.



# Default Deny

## **APEX Security Model**

- ✓ Unregistered Servers DENIED
- ✓ Unapproved Tools DENIED
- ✓ NO direct credentials
- ✓ Managed identity ONLY
- ✓ Audit MANDATORY
- ✓ Least privilege ENFORCED

# Rogue MCP Server

## UNREGISTERED MCP SERVER



```
Tool call: read_customer_database
params: { "query": "SELECT * FROM Customers;
DROP TABLE AuditEvents;" }
origin: mcp://unverified-7f3a.external
subscription: (none)
```

## APEX-GOVERNED RESOURCE

APEX

```
>>> APEX policy check

layer 1 APIM: FAIL no subscription key
layer 2 policy: FAIL source not registered
layer 3 identity: SKIPPED
layer 4 RBAC: SKIPPED
layer 5 audit: LOGGED
```

-----

Status: DENIED

Reason: Unregistered AI capability source

# Approved Vs Rogue Server

# Approved MCP Server

## REGISTERED MCP SERVER



```
Tool call: customer_lookup
params: { "customer_id": 1842 }
origin: mcp://corp-crm-agent.internal
subscription: apim-key-9c4e (valid)
```

## APEX-GOVERNED RESOURCE

APEX

```
>>> APEX policy check

layer 1 APIM: PASS key + JWT valid
layer 2 policy: PASS ReadOnlyMediumRisk
layer 3 identity: PASS managed identity
layer 4 RBAC: PASS db_datareader
layer 5 audit: LOGGED
```

-----

Status: Approved

Policy: ReadOnlyMediumRisk

# Entitlement Matrix

| User        | SQL    | Blob   | Outlook           | Web                 |
|-------------|--------|--------|-------------------|---------------------|
| FinanceBot  | read   | read   | internal-only     | none                |
| SalesAgent  | read   | none   | external<br>email | approved<br>domains |
| RogueServer | denied | denied | denied            | denied              |

□ LIVE EVIDENCE

## APEX in action

Working software, not slideware. Each surface below maps to a deployed APEX capability.

### 1 Tool registration

apex.console

Register new MCP server

name: corp-crm-agent

origin: mcp://corp-crm-agent.internal

owner: data-platform-team

tools: customer\_lookup, order\_status

→ Submitted for approval

Every MCP server entering the enterprise is named, owned, and scoped before it can call anything.

### 2 Approval workflow

queue.apex

corp-crm-agent

PENDING REVIEW

Risk tier: Medium

Data scope: Customers (read-only)

Approver: Security Officer

Approve

Deny

Hold

Human-in-the-loop approval by Security before any MCP server reaches production data.

### 3 Audit trail

log.apex

14:22:01 PASS corp-crm-agent →  
customer\_lookup(1842)

14:22:09 PASS corp-crm-agent → order\_status(7821)

14:22:14 DENY unverified-7f3a → read\_db(\*)

14:22:18 PASS finance-bot → ledger\_query(Q3)

14:22:25 FLAG corp-crm-agent → bulk\_export

Immutable record of every AI tool call with identity, policy result, and full payload.

### 4 Policy enforcement

policy.apex

ReadOnlyMediumRisk

Read operations

ALLOW

Write operations

DENY

Bulk export (> 1000 rows)

REQUIRE APPROVAL

PII fields

MASK

Rate limit

100 /  
min

Granular policy tiers map MCP servers to allowed operations, data scope, and quotas.

# APEX In Action

# Screenshots

## Register Tool

Tool Name

customer\_lookup

Display Name

Customer Lookup

Description

Lookup customer by ID through governed SQL endpoint

Backend Path

/tool/sql/customer/{customer\_id}

Resource Type

Azure SQL

Resource Name

Customers

Access Type

read-only

Risk Level

medium

Register Tool

## Registered Tools / Approval Flow

Refresh Tools

| ID | Tool            | Display         | Backend                          | Resource             | Risk   | Status  | Action  |
|----|-----------------|-----------------|----------------------------------|----------------------|--------|---------|---------|
| 11 | customer_lookup | Customer Lookup | /tool/sql/customer/{customer_id} | Azure SQL: Customers | medium | pending | Approve |

## Create Policy Template

Policy Name

ReadOnlyMediumRiskCustom

Risk Level

medium

Rate Limit Per Minute

30

Requires JWT

☒

Requires Audit

☒

Allow After Hours

☒

Allow Writes

☐

Allow Bulk Export

☐

Create Policy Template

## Policy Templates

Refresh Policy Templates

| ID | Name                     | Risk   | JWT  | Audit | Rate Limit | Writes | Bulk Export |
|----|--------------------------|--------|------|-------|------------|--------|-------------|
| 1  | ReadOnlyLowRisk          | low    | true | true  | 60         | false  | false       |
| 2  | ReadOnlyMediumRisk       | medium | true | true  | 30         | false  | false       |
| 3  | HighRiskRestricted       | high   | true | true  | 10         | false  | false       |
| 4  | ReadOnlyMediumRiskCustom | medium | true | true  | 30         | false  | false       |

## Audit Trail

```
{
  "count": 20,
  "events": [
    {
      "audit_event_id": 20,
      "timestamp_utc": "2026-05-16 19:03:03.550270",
      "tool_name": "read_blob_document",
      "resource_accessed": "Azure Blob: documents/demo.txt",
      "access_model": "managed_identity_governed_blob_endpoint",
      "status": "success"
    },
    {
      "audit_event_id": 19,
      "timestamp_utc": "2026-05-16 19:02:56.423734",
      "tool_name": "sql_customer_lookup",
      "resource_accessed": "Azure SQL: Customers/1",
      "access_model": "managed_identity_governed_sql_endpoint",
      "status": "success"
    },
    {
      "audit_event_id": 18,
      "timestamp_utc": "2026-05-16 08:24:09.811893",
      "tool_name": "create_policy_template",
      "resource_accessed": "Policy Template: ReadOnlyMediumRiskCustom",
      "access_model": "admin_policy_workflow",
      "status": "created"
    },
    {
      "audit_event_id": 17,
      "timestamp_utc": "2026-05-16 08:22:25.999711",
      "tool_name": "approve_tool",
      "resource_accessed": "Tool Registry: ToolId 1",
      "access_model": "admin_approval_workflow",
      "status": "approved"
    }
  ]
}
```



# Enterprise Value

## Enable AI without enabling risk

Enterprise AI fails on governance, not on models. APEX is the governance layer.

### □ Prevent AI data breaches

Unregistered MCP servers and rogue agents cannot reach enterprise data.

Default-deny on every capability source

### □ Satisfy AI audit requirements

Every tool call, identity, and policy decision is logged and queryable.

SOC 2, HIPAA, SOX, EU AI Act ready

### □ Ship AI capabilities faster

Approved MCP servers deploy through a registration workflow, not security review backlog.

Weeks to days for new AI tooling

### □ Control AI consumption

Per-source quotas, rate limits, and policy tiers prevent runaway LLM and data costs.

Spend visibility per agent, tool, and policy

### The enterprise position

Without governance, AI is either blocked by security or deployed with hidden risk. APEX is the layer that lets the enterprise say yes.



# Contact

## APEX

Enterprise AI without governance is enterprise risk without controls.

APEX is the layer between AI capability and enterprise data.

LIVE DEMO

[opschainai.com](https://opschainai.com)

Working APEX reference architecture, governing MCP servers against real enterprise data sources.



### Live demo

See APEX deny a rogue MCP server in real time against a working enterprise data source.



### Pilot scoping

Two-week engagement to map your AI capability surface and define a governance baseline.



### Technical briefing

Deep dive on the architecture, threat model, and governance layers for your security team.

### CONTACT

#### Paul Hasselbring

Founder & Principal Architect, NPM Technologies

- ✉ [paulh@npmit.com](mailto:paulh@npmit.com)
- ☎ (855) 676-8324
- 🌐 [linkedin.com/in/paulhasselbring](https://linkedin.com/in/paulhasselbring)

### DEPLOYED ON

Azure

Entra ID

APIM

Container Apps